






A-Fig3 : Our bounded inference cost vs. history. Left: Conditioning memory tokens saturate at 512 after chunk 2, as the fixed-size EM-compressed bank merges older blocks on overflow. Middle: Per-chunk latency stays flat at ~ 74 s (solid: measured on clean GPU; open: fixed-token plateau). Right: Memory-chunk occupancy caps at 2. Cost is independent of total video length.